

# Process Synchronization

Mehdi Kargahi  
School of ECE  
University of Tehran  
Summer 2016

# Producer-Consumer (Bounded Buffer)

## ■ Producer:

```
while (true) {  
    /* produce an item in next_produced */  
  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

## ■ Consumer:

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
  
    /* consume the item in next_consumed */  
}
```

# Race Condition

## ■ Producer

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

## ■ Consumer

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

$T_0$ :	producer	execute	$register_1 = counter$	{ $register_1 = 5$ }
$T_1$ :	producer	execute	$register_1 = register_1 + 1$	{ $register_1 = 6$ }
$T_2$ :	consumer	execute	$register_2 = counter$	{ $register_2 = 5$ }
$T_3$ :	consumer	execute	$register_2 = register_2 - 1$	{ $register_2 = 4$ }
$T_4$ :	producer	execute	$counter = register_1$	{ $counter = 6$ }
$T_5$ :	consumer	execute	$counter = register_2$	{ $counter = 4$ }

# When CPU Scheduling Occurs

---

1. Running process switches to wait state
  - I/O request
  - Waiting for termination of one of the child processes
2. Running process switches to ready state (timer interrupt)
3. Waiting process switches to ready state (I/O completion)
4. Running process terminates

Under non-preemptive scheduling: 1 and 4 are legal scheduling points

Under preemptive scheduling: 1 through 4 are legal Scheduling points

# Critical Sections

## ■ Structure of a Typical Process

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (TRUE);
```

## ■ Requirements for Solving critical-section problem:

- Mutual exclusion
- Progress
- Bounded waiting

# Two-Process Solutions

## ■ Algorithm 1

■ Shared: int turn;

$P_i$ :

```
while (TRUE) {  
    while (turn!=i);  
    critical-section  
    turn = 1-i;  
    remainder-section  
}
```

Mutual exclusion yes

Progress no

Bounded waiting yes

# Two-Process Solutions

## ■ Algorithm 2

■ Shared: boolean flag[2];

$P_i$ :

flag[i] = FALSE;

do {

    flag[i] = TRUE;

    while (flag[j]);

    critical-section

    flag[i] = FALSE;

    remainder-section

} while (TRUE);

Mutual exclusion yes

Progress no

Bounded waiting yes

# Peterson's Solution

```
do {
```

```
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);
```

critical section

```
    flag[i] = FALSE;
```

remainder section

```
} while (TRUE);
```



# Eisenberg & McGuire Solution

---

```
enum pstate {idle, want_in, in_cs};  
pstate flag[n];  
int turn;
```

```

do {
    while (TRUE) {
        flag[i] = want_in;
        j = turn;

        while (j != i) {
            if (flag[j] != idle) {
                j = turn;
            }
            else
                j = (j + 1) % n;
        }

        flag[i] = in_cs;
        j = 0;

        while ( (j < n) && (j == i || flag[j] != in_cs) )
            j++;

        if ( (j >= n) && (turn == i || flag[turn] == idle) )
            break;
    }

    // critical section

    j = (turn + 1) % n;

    while (flag[j] == idle)
        j = (j + 1) % n;

    turn = j;
    flag[i] = idle;

    // remainder section
}while (TRUE);

```

# Synchronization Hardware

## ■ Atomic instructions: TestAndSet

```
boolean TestAndSet(boolean *target) {  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

```
do {  
    while (TestAndSetLock(&lock))  
        ; // do nothing  
  
    // critical section  
  
    lock = FALSE;  
  
    // remainder section  
}while (TRUE);
```

# Synchronization Hardware

## ■ Atomic instructions: Swap

```
void Swap(boolean *a, boolean *b) {  
    boolean temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

```
do {  
    key = TRUE;  
    while (key == TRUE)  
        Swap(&lock, &key);  
  
    // critical section  
  
    lock = FALSE;  
  
    // remainder section  
}while (TRUE);
```

# Synchronization Hardware

---

## ■ Atomic instructions: Compare\_and\_Swap

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
} while (true);
```

# Synchronization Hardware

---

- Spin lock
  - Methods that use busy waiting
- n-process solution with bounded-waiting

```
boolean waiting[n];  
boolean lock;
```

# Synchronization Hardware

```
do {
    waiting[i] = TRUE;
    key = TRUE;
    while (waiting[i] && key)
        key = TestAndSet(&lock);
    waiting[i] = FALSE;

    // critical section

    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;

    if (j == i)
        lock = FALSE;
    else
        waiting[j] = FALSE;

    // remainder section
}while (TRUE);
```

# Mutex Locks

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}
```

```
release() {  
    available = true;  
}
```

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```



# Semaphores

- A semaphore S is an integer variable
  - Two atomic operations

```
wait(S) {  
    while S <= 0  
        ; // no-op  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

# Semaphore Usage

## ■ Critical-section

```
do {  
    waiting(mutex);  
  
    // critical section  
  
    signal(mutex);  
  
    // remainder section  
}while (TRUE);
```

## ■ Process synchronization (Counting vs. Binary Semaphore)

```
S1;  
signal(synch);
```

```
wait(synch);  
S2;
```

## ■ *Problems*: Spin-lock & bounded-waiting

# A Better Definition

```
typedef struct {
    int value;
    struct process *list;
} semaphore;

wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

# Semantics

---

- The value of semaphore specifies the number of processes waiting on that semaphore
- It is assumed that wait and signal are executed atomically, while this is not simply valid
- *Method 1*: disabling interrupts to prevent interleaving of wait and signal operations on a semaphore
  - In multiprocessor systems, interrupts on all processors should be disabled resulting in quite low performance
- *Method 2*: using one of the previous critical-section methods (busy waiting) → More proper for MP systems
  - Wait and signal have a few instructions resulting in short critical sections → Busy waiting is not problematic

# Deadlocks

$P_0$	$P_1$
<code>wait(S);</code>	<code>wait(Q);</code>
<code>wait(Q);</code>	<code>wait(S);</code>
<code>.</code>	<code>.</code>
<code>.</code>	<code>.</code>
<code>.</code>	<code>.</code>
<code>signal(S);</code>	<code>signal(Q);</code>
<code>signal(Q);</code>	<code>signal(S);</code>

- What is the difference between deadlock and starvation?

# Priority Inversion

---

- PI occurs in systems with more than two priorities. So one solution is to have at most two priority levels.
- PIP: Priority-Inheritance Protocol
- PI and the Mars Pathfinder

# Classical Problems of Synchronization

## ■ The bounded-buffer problem

### Producer

```
do {  
    . . .  
    // produce an item in nextp  
    . . .  
    wait(empty);  
    wait(mutex);  
    . . .  
    // add nextp to buffer  
    . . .  
    signal(mutex);  
    signal(full);  
}while (TRUE);
```

### Consumer

```
do {  
    wait(full);  
    wait(mutex);  
    . . .  
    // remove an item from buffer to nextc  
    . . .  
    signal(mutex);  
    signal(empty);  
    . . .  
    // consume the item in nextc  
    . . .  
}while (TRUE);
```

# The Readers-Writers Problem

Writer

```
do {  
    wait(wrt);  
    . . .  
    // writing is performed  
    . . .  
    signal(wrt);  
}while (TRUE);
```

```
semaphore mutex, wrt;  
int readcount;
```


Reader

```
do {  
    wait(mutex);  
    readcount++;  
    if (readcount == 1)  
        wait(wrt);  
    signal(mutex);  
    . . .  
    // reading is performed  
    . . .  
    wait(mutex);  
    readcount--;  
    if (readcount == 0)  
        signal(wrt);  
    signal(mutex);  
}while (TRUE);
```



# The Dining Philosophers Problem

---



# The Dining Philosophers Problem

```
semaphore chopstick[5];  
  
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1) % 5]);  
    . . .  
    // eat  
    . . .  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    . . .  
    // think  
    . . .  
}while (TRUE);
```

- Problem?
- Three solutions !!
- Is a deadlock-free solution also starvation-free?

# Monitors

---

- Problems with semaphores (programming faults):

```
signal(mutex);  
...  
critical section  
...  
wait(mutex);
```

```
wait(mutex);  
...  
critical section  
...  
wait(mutex);
```

# Syntax of a Monitor

```
monitor monitor name
{
    // shared variable declarations


    procedure P1 ( . . . ) {
        . . .
    }

    procedure P2 ( . . . ) {
        . . .
    }

    .
    .
    .
    procedure Pn ( . . . ) {
        . . .
    }

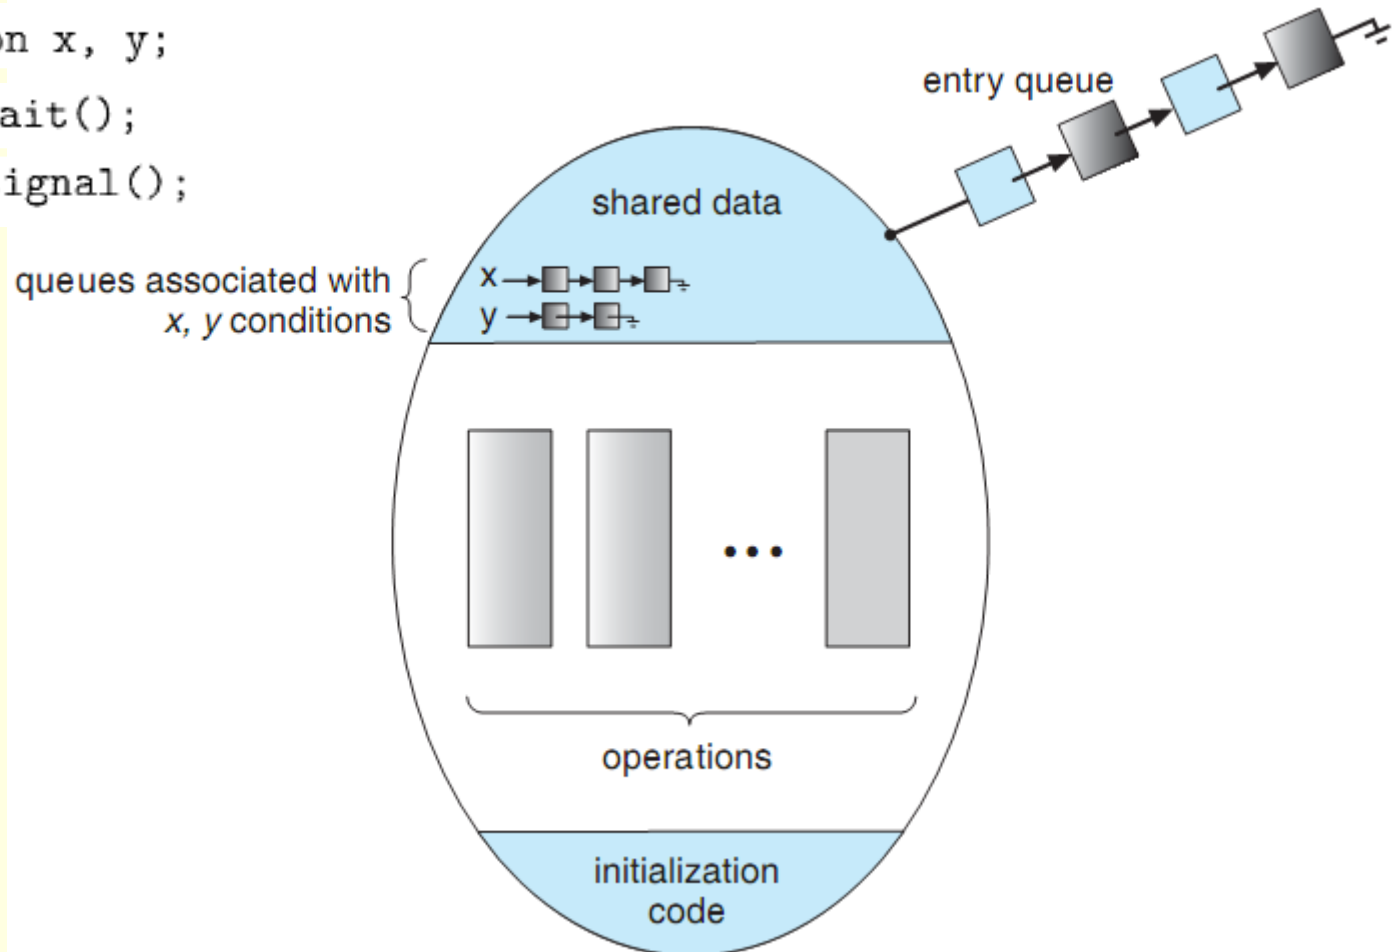
    initialization code ( . . . ) {
        . . .
    }
}
```

# A Monitor Scheme



# Monitor with Condition Variables

```
condition x, y;  
x.wait();  
x.signal();
```



# When `x.signal()` is Invoked by *P*

---

- Two possibilities exist:

1. **Signal and wait.** *P* either waits until *Q* leaves the monitor or waits for another condition.
2. **Signal and continue.** *Q* either waits until *P* leaves the monitor or waits for another condition.

# Dining-Philosophers Solution using Monitors

---

```
enum {thinking, hungry, eating} state[5];
```

```
condition self[5];
```

```
dp.pickup(i);  
...  
eat  
...  
dp.putdown(i);
```



# Dining-Philosophers Solution using Monitors

```
monitor dp
```

```
{  
    enum {THINKING, HUNGRY, EATING}state[5];  
    condition self[5];
```

```
void pickup(int i) {  
    state[i] = HUNGRY;  
    test(i);  
    if (state[i] != EATING)  
        self[i].wait();  
}
```

```
void putdown(int i) {  
    state[i] = THINKING;  
    test((i + 4) % 5);  
    test((i + 1) % 5);  
}
```

```
void test(int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING)) {  
        state[i] = EATING;  
        self[i].signal();  
    }  
}
```

```
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}
```

# Implementing a Monitor using Semaphores

- External procedure F is replaced by

```
wait(mutex);  
    ...  
    body of F  
    ...  
if (next_count > 0)  
    signal(next);  
else  
    signal(mutex);
```

## x.wait()

```
x_count++;  
if (next_count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x_sem);  
x_count--;
```

## x.signal()

```
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```

# Resuming Processes within a Monitor

- `x.wait(c)`, where `c` is called a priority number

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;

    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }

    void release() {
        busy = FALSE;
        x.signal();
    }

    initialization_code() {
        busy = FALSE;
    }
}
```

```
R.acquire(t);
...
    access the resource;
...
R.release();
```

# Similar Mechanisms

---

- Down: wait
- Up: signal

# Problems with Monitors

---

- A process might access a resource without first gaining access permission
- A process might never release a resource
- A process might attempt to release a resource that it never requested
- A process might request the same resource twice

# Atomic Transactions

---

- Transaction
  - A collection of instructions that performs a single logical function
- A transaction may **commit** or **abort**
- When a transaction is aborted, the state of the data accessed by it is restored to what it was before executing that transaction, i.e., it is **rolled back**
- The execution of a transaction should be assumed as **atomic**

# Different Types of Storage

---

- Volatile storage
- Non-Volatile storage
- Stable storage

# Log-Based Recovery

---

- Log is stored on the stable storage
- Write-ahead log
  - Every log precedes a single write operation by
    - Transaction name
    - Data item name
    - Old value
    - New value
  - Special log records
    - $\langle T_i \text{ starts} \rangle$
    - $\langle T_i \text{ commits} \rangle$



# Recovery Algorithm

---

- $Undo(T_i)$ : restores the value of all data updates by  $T_i$  to the old values
- $Redo(T_i)$ : sets the value of all data updates by  $T_i$  to the new values
- If  $T_i$  aborts, we need to execute  $undo(T_i)$
- After a system failure
  - Transaction  $T_i$  needs to be undone if the log contains the  $\langle T_i \text{ starts} \rangle$  record but does not contain the  $\langle T_i \text{ commits} \rangle$  record.
  - Transaction  $T_i$  needs to be redone if the log contains both the  $\langle T_i \text{ starts} \rangle$  and the  $\langle T_i \text{ commits} \rangle$  records.

# Concurrent Atomic Transactions

- Serializability
  - Restrictive
- Serial vs. non-serial schedules

## Serial

$T_0$	$T_1$
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)

## Non-serial but serializable

$T_0$	$T_1$
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

- Conflicting operations:  $O_i$  and  $O_j$  are conflicting if they access the same data item and at least one of them is a *write* operation

# Concurrent Atomic Transactions

---

- A conflict serializable schedule
  - If a schedule  $S$  can be transformed to a serial schedule  $S'$  by a series of swaps of non-conflicting operations

# Locking Protocol

---

- Shared vs. exclusive lock
- Two-phase locking (2PL) protocol
  - **Growing phase.** A transaction may obtain locks but may not release any lock.
  - **Shrinking phase.** A transaction may release locks but may not obtain any new locks.
- Properties
  - Ensures conflict Serializability
  - It is not deadlock-free
  - There may be conflict-serializable schedules that cannot be obtained by 2PL
    - Additional information on transactions may improve 2PL

# Timestamp-Based Protocols

## ■ Each transaction $T_i$ is assigned a timestamp in an ascending order of time by the system

- Use the value of the system clock as the timestamp; that is, a transaction's timestamp is equal to the value of the clock when the transaction enters the system. This method will not work for transactions that occur on separate systems or for processors that do not share a clock.
- Use a logical counter as the timestamp; that is, a transaction's timestamp is equal to the value of the counter when the transaction enters the system. The counter is incremented after a new timestamp is assigned.

## ■ Each data item $Q$ has two timestamps

**W-timestamp**( $Q$ ) denotes the largest timestamp of any transaction that successfully executed  $\text{write}(Q)$ .

**R-timestamp**( $Q$ ) denotes the largest timestamp of any transaction that successfully executed  $\text{read}(Q)$ .

# Timestamp-Based Protocols

## ■ If transaction $T_i$ issues $read(Q)$

- If  $TS(T_i) < W\text{-timestamp}()$ , then  $T_i$  needs to read a value of  $Q$  that was already overwritten. Hence, the read operation is rejected, and  $T_i$  is rolled back.
- If  $TS(T_i) \geq W\text{-timestamp}(Q)$ , then the read operation is executed, and  $R\text{-timestamp}(Q)$  is set to the maximum of  $R\text{-timestamp}(Q)$  and  $TS(T_i)$ .

## ■ If transaction $T_i$ issues $write(Q)$

- If  $TS(T_i) < R\text{-timestamp}(Q)$ , then the value of  $Q$  that  $T_i$  is producing was needed previously and  $T_i$  assumed that this value would never be produced. Hence, the write operation is rejected, and  $T_i$  is rolled back.
- If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of  $Q$ . Hence, this write operation is rejected, and  $T_i$  is rolled back.
- Otherwise, the write operation is executed.

## ■ Each transaction that is rolled is restarted with a new timestamp

# Timestamp-Based Protocols

---

- Properties

- Deadlock-free: no transaction ever waits
- Not stronger than 2PL and vice versa

# Approaches for Multicore Processors

---

- Problems if we use classic locks
  - Increased risk of race condition
  - Increased risk of deadlock
  - Low scalability because of contention between very large number of threads
- Solutions that better scale for large number of cores to support designing thread-safe concurrent applications
  - In programming languages
  - In hardware