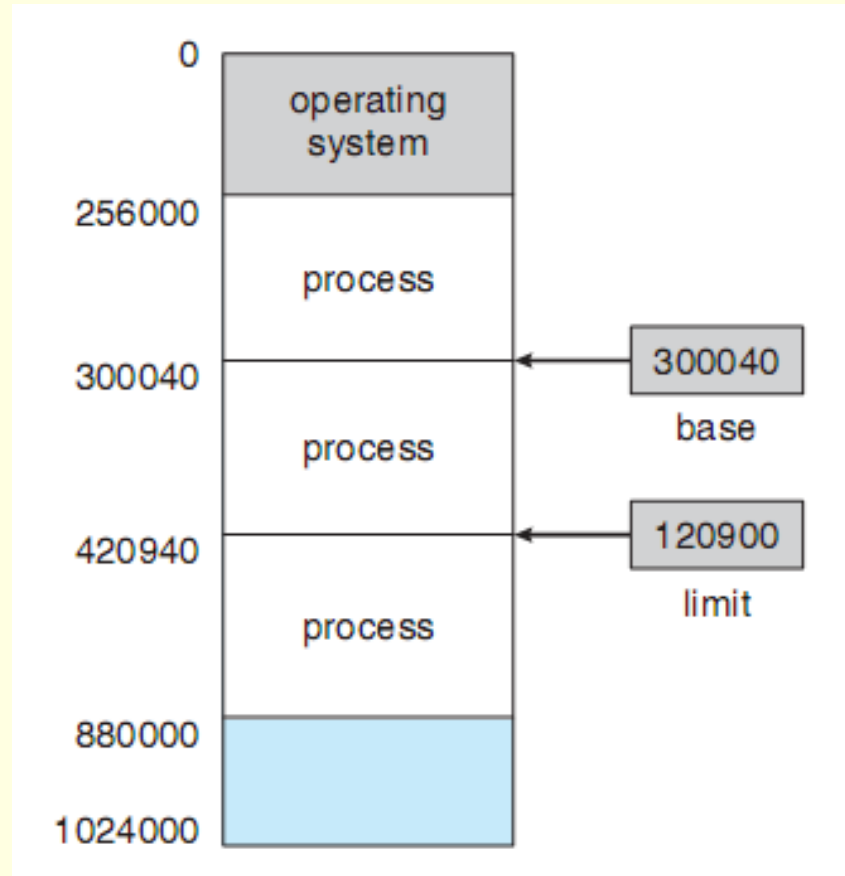


Memory Management (Main Memory)

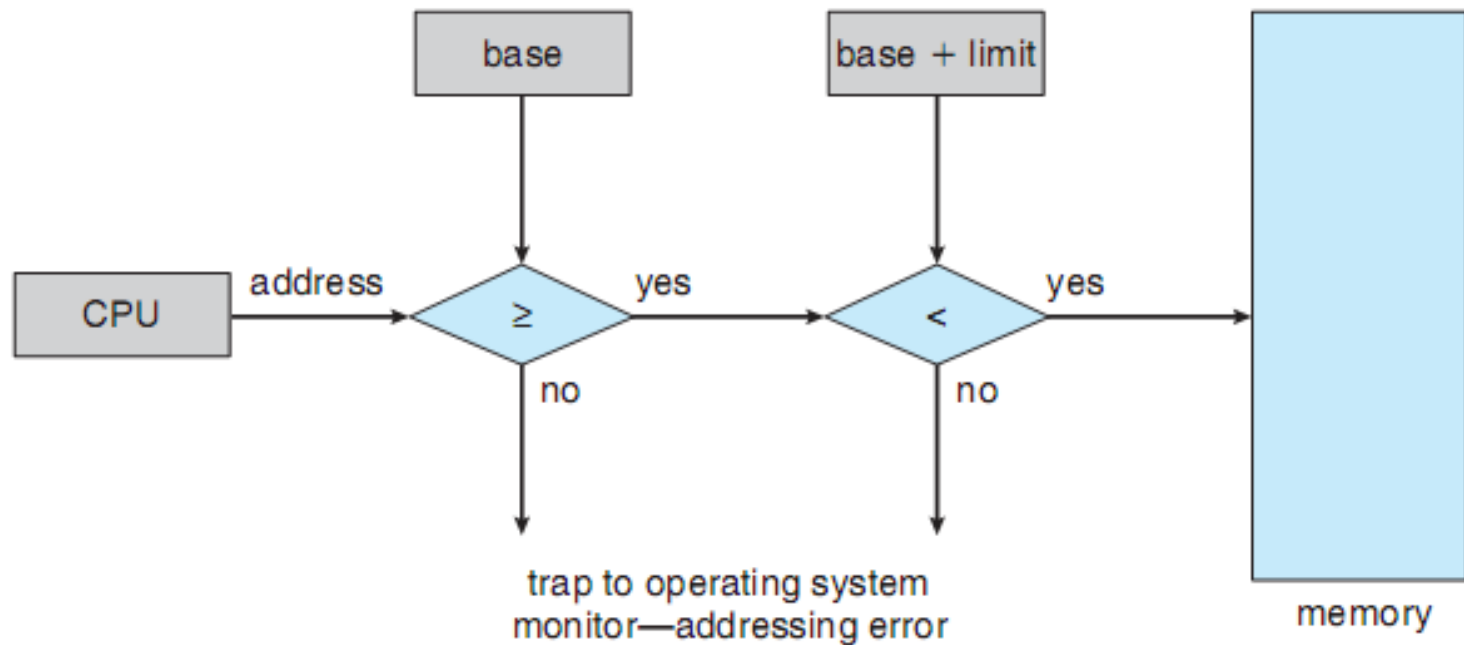
Mehdi Kargahi
School of ECE
University of Tehran
Summer 2016

Hardware Address Protection



Hardware Address Protection

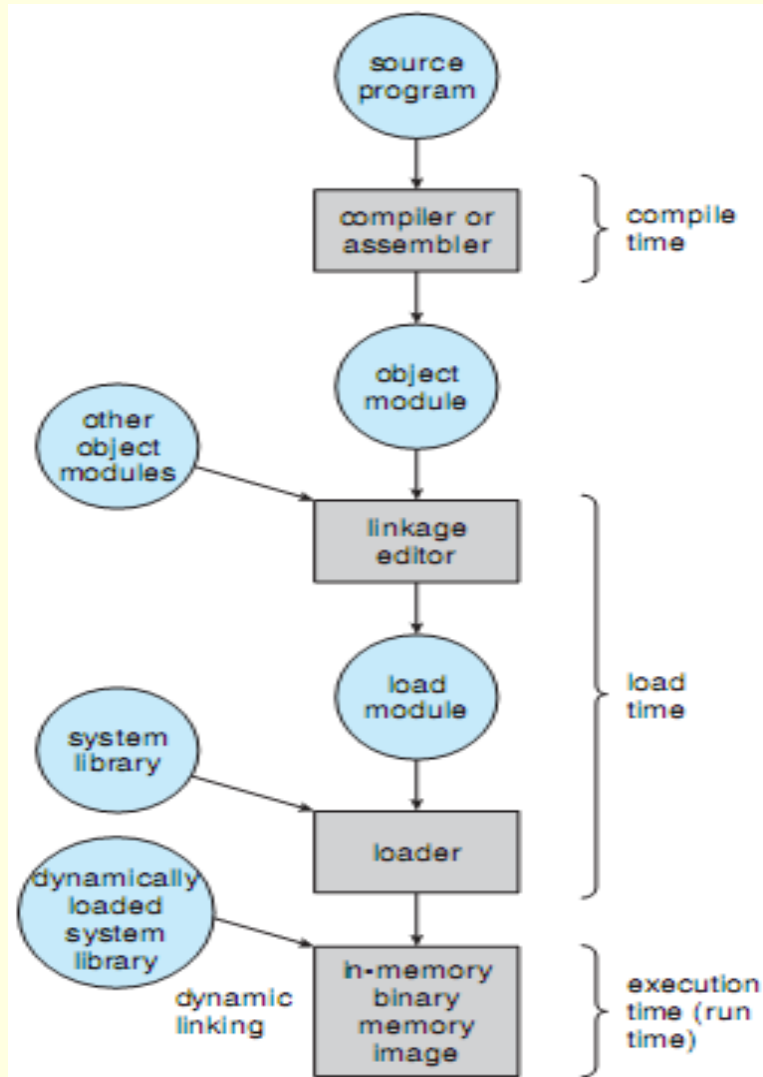
■ Base and Limit Registers



Address Binding

- Binding: a mapping from one address space to another
 - Compiler time (absolute code)
 - Load time (relocatable code)
 - Execution time (process can be moved during its execution)
 - Code may partially be loaded
 - Available in most general-purpose OS

Address Binding



Logical versus Physical Address Space

- Logical address: the address generated by the CPU
- Physical address: the address seen by the memory unit (loaded into memory address register)
- Virtual address: logical address in execution time binding which is different from the respective physical address
 - Mapping is done using MMU (memory-management unit)

Dynamic Loading

- A routine is not loaded until it is called
- For better memory space utilization
- Process size is not limited to the size of physical memory
- An almost general rule: 10% main program and 90% for exception handling
- Dynamic loading does not require special support from the OS, but OS may help the programmer by providing special library routines

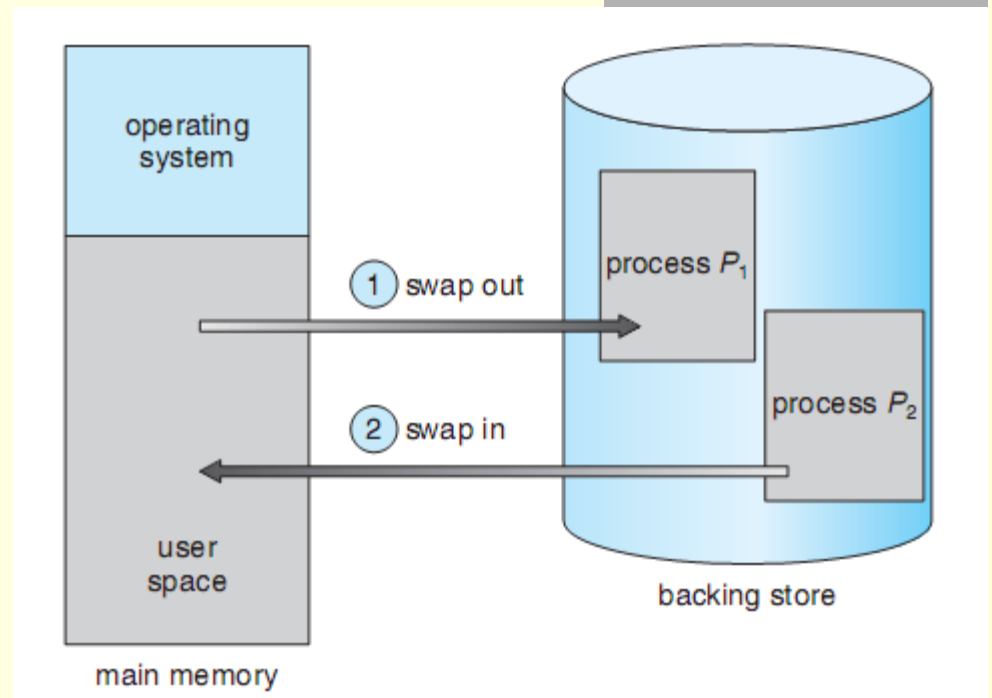
Overlay Technique

- A weak memory management technique without using the OS
- Ex.
 - An assembler
 - Pass1 (70K) & symbol table (20K)
 - Pass2 (80K) & common routines (30K)
 - 200K is needed, while 150K is available
 - $ST (20K) + CR (30K) + OV (10K) + Pass1 \ \& \ Pass2 (80K) = 140K$

Dynamic Linking and Shared Libraries

- Static linking versus dynamic linking
- Here, linking rather than loading is postponed until execution time
- To save both disk space and memory
- Stub: a small piece of code indicating how to locate the appropriate library routine
- Only one copy of the library routine is loaded
- Library bug fixes are much more simpler
 - Shared libraries: related to different library versions for different processes → previously compiled programs are not affected by the new versions
- Needs support from the OS for usage of a library by multiple processes

Swapping



■ Ex.

- Average latency and seek time: 8ms
- Transfer rate to HDD: 40 MB/s
- User process 10 MB
- $10/40 + 0.008 = 258 \text{ ms} \rightarrow 2 \times 258 = 516 \text{ ms}$

Swapping

- What if any pending I/O
 - Solution 1: Never swap a process with pending I/O
 - Solution 2: Execute I/O operations only in the OS buffers

Swapping on Mobile Systems

- Mobile devices generally use flash memory rather than more spacious hard disks as their persistent storage.
- Why mobile operating-system designers avoid swapping:
 - Space constraints
 - The limited number of writes that flash memory can tolerate before it becomes unreliable
 - The poor throughput between main memory and flash memory in these devices.

Contiguous Memory Allocation

- OS usually resides in low memory space rather than high memory space because of the location of the interrupt vector
- *Memory manager* should be aware of current *holes* and *used pieces* of memory and do the allocation
- Fixed-size partitions
- *First-fit*, *best-fit*, *worst-fit*, and *next-fit* strategies

Fragmentation

- External fragmentation

- Ex.: Memory size: 2560K , RR ($q=1$)
- OS (0-400K)
- P1 (600K, 10ms), P2(1000K, 5ms), P3(300K, 20ms), P4(700K, 8ms), P5 (500K, 15ms)

- Internal fragmentation

- Ex.: 18464 bytes are free, but 18462 bytes are needed → a good allocation policy may use all the free memory block to prevent some overheads

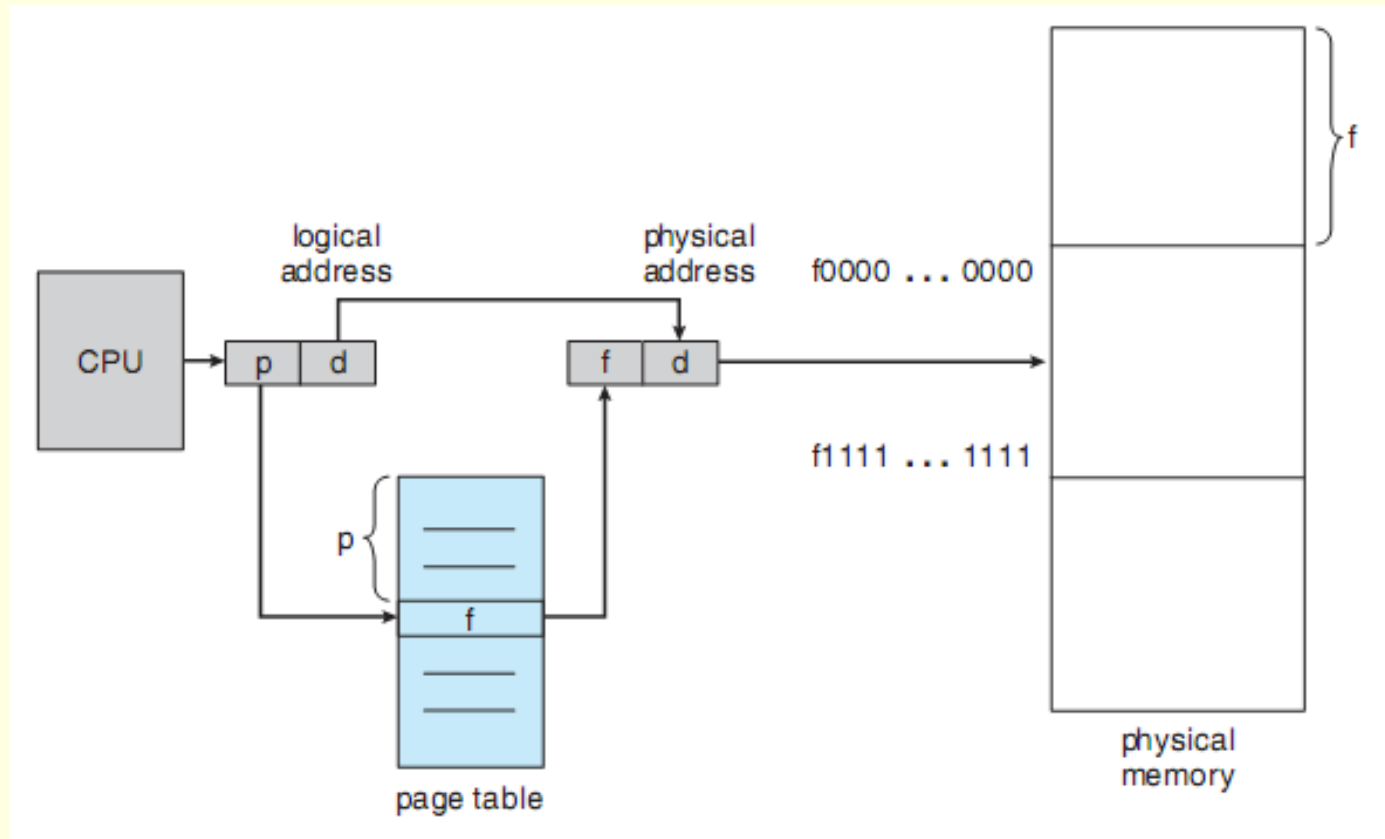
Solutions to External Fragmentation

- Compaction
 - Only applicable if relocation is dynamic and is done at execution time
 - Trying on the previous example!
- Using non-contiguous logical address space
 - Paging
 - Segmentation

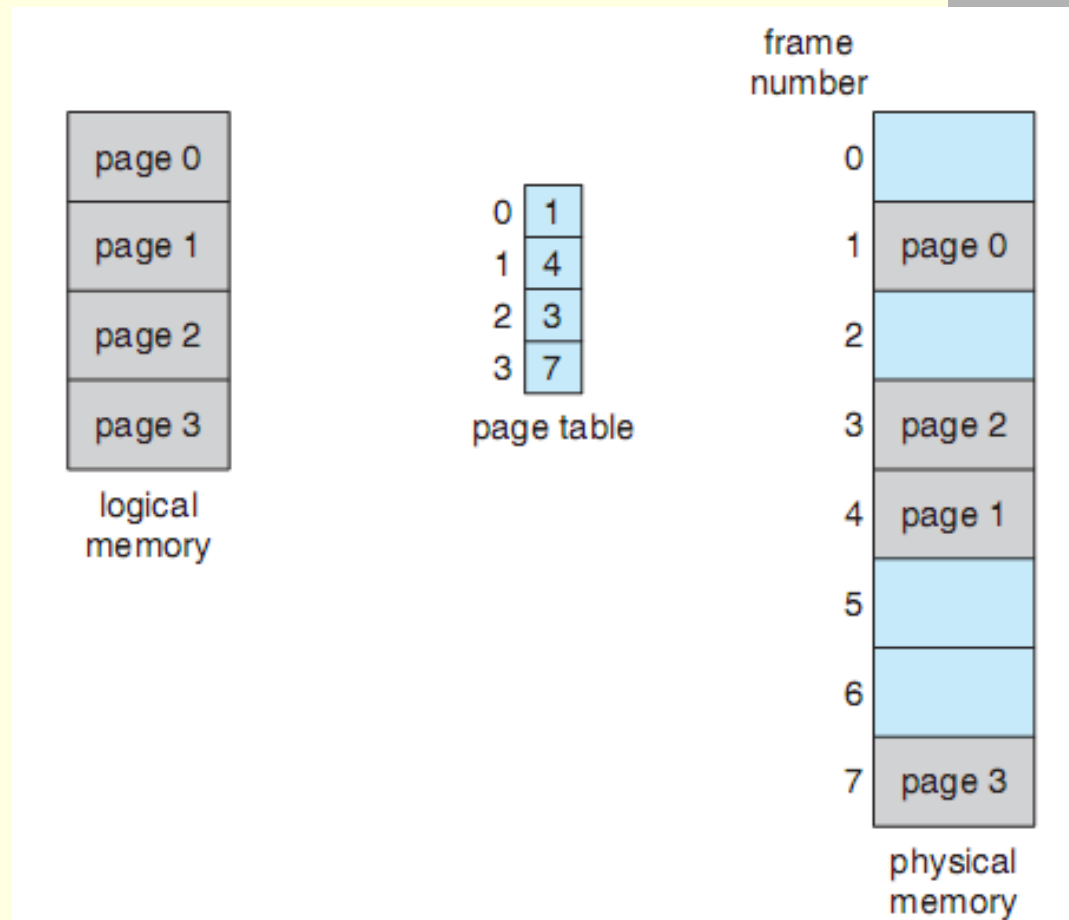
Paging

- Basic method
 - Fixed size frames and pages
- Advantages
 - Contiguous space is not required
 - Fitting memory chunks onto backing store is not problematic
 - Fitting pages onto frames is straightforward

Hardware Support for Paging



Example



Calculating Physical Address

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory


0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	


physical memory

Frame Table



Logical Address

- Page size and frame size are defined by hardware
- Some operating systems support variable page sizes
- Ex.
 - Size of logical address space: 2^m
 - Page size: 2^n words



Where the Page Table is Stored?

- OS maintains a copy of the page table
 - Whenever a process makes a system call and passes a pointer to a buffer as a parameter, binding should be done properly
 - CPU dispatcher also uses this copy to define the hardware page table when a process is to be allocated the CPU

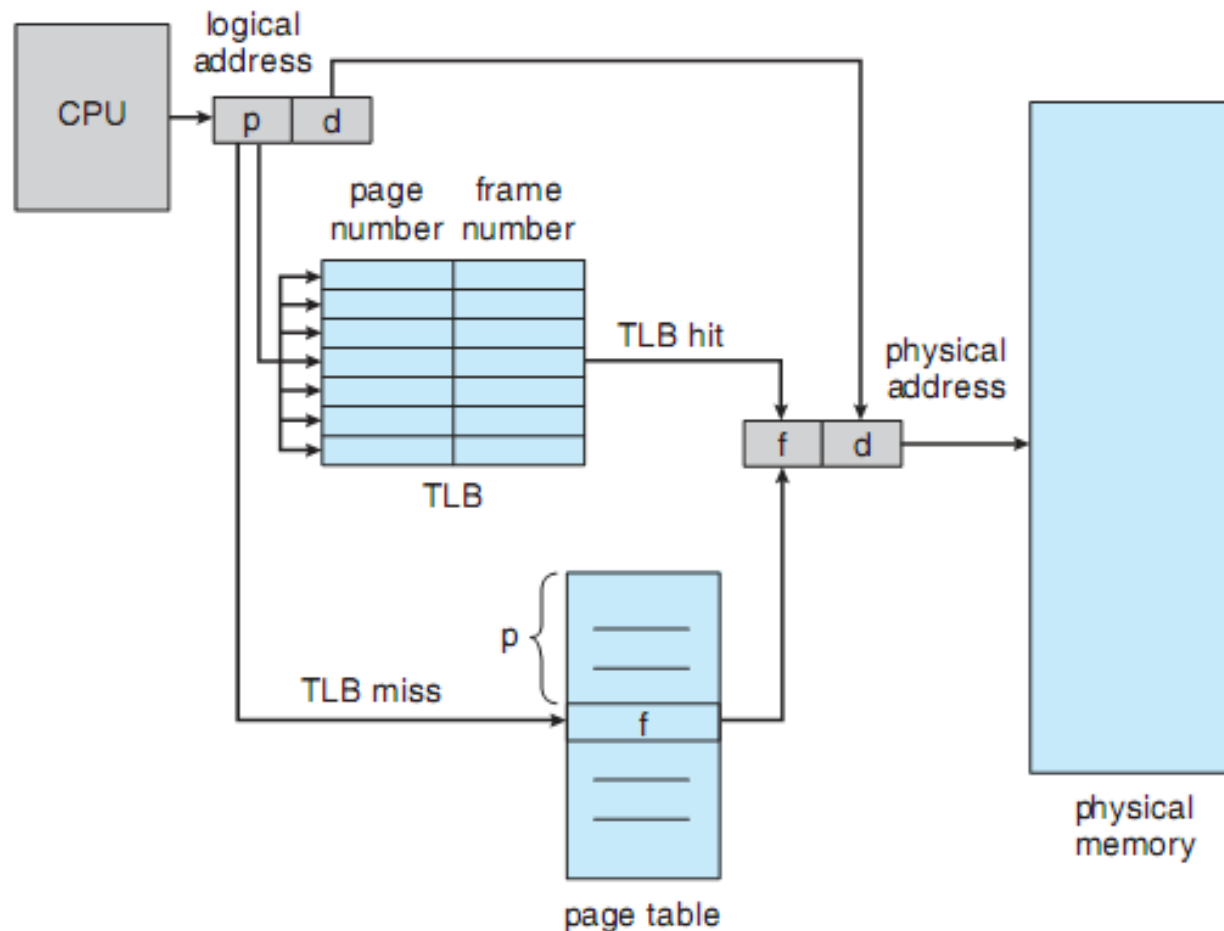
Hardware Support

- Where the page table is stored?
 - In special CPU registers
 - Modification is privileged
 - Limitation in size
 - Relatively high context-switch time
 - In memory (having a PTBR)
 - Modifying PTBR is privileged
 - Lower context-switch time
 - Memory access time is larger

Hardware Support

- Using Translation Look-aside Buffer (TLB)
 - TLB: a special small fast lookup hardware cache (associative high speed memory)
 - Number of entries in TLB are typically between 64 and 1024
 - What should be done if a TLB miss occurs?
 - TLB full → Using a replacement policy
 - **Wired down** TLBs (non-removable) are used for kernel code

Paging Hardware with TLB



Hardware Support

- Some TLBs store address-space identifiers (ASIDs)
- ASID: uniquely identifies a process
- By checking the ASID for each virtual page number, address space protection is done. Otherwise the TLB should be flushed on each CS.
- Effective memory access time
 - Ex.: Hit ratio: 80% (80386 with 1 TLB), TLB access: 20^{ns} , memory access: 100^{ns}
 - Effective access time: $80\% * 120 + 20\% * 220 = 140^{\text{ns}}$
 - hr= 98% (80486 with 32 TLBs) → EAT: 122^{ns}

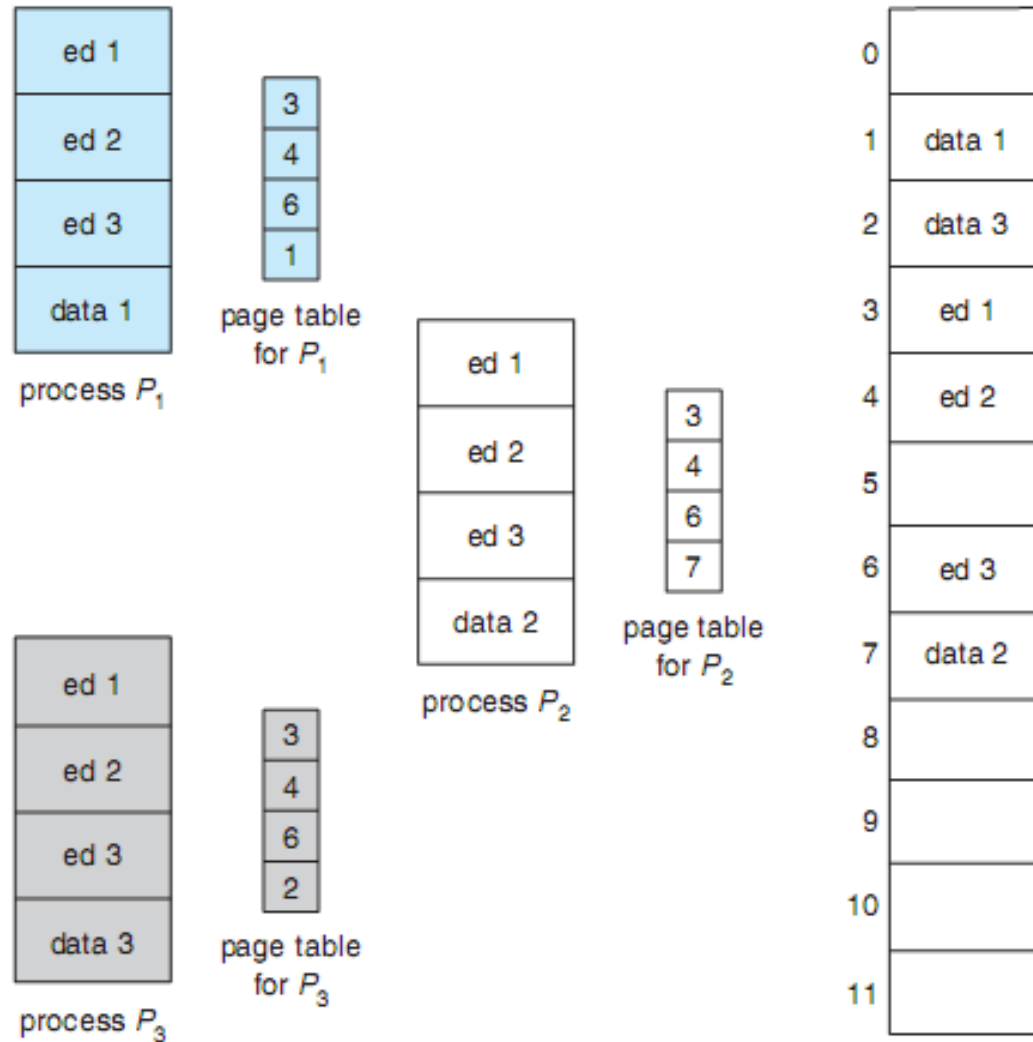
-
- The diagram illustrates the mapping of virtual memory to physical memory using a page table. It shows three main components:
- Virtual Address Space:** A vertical stack of pages labeled "page 0" through "page 5". The address range "00000" is associated with "page 0", and "10,468" and "12,287" are associated with "page 5".
 - Page Table:** A table mapping virtual page numbers to physical frame numbers and a valid-invalid bit.

Virtual Page	Frame Number	Valid-Invalid Bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i
 - Physical Memory:** A vertical stack of frames labeled "0" through "n". Pages are mapped to frames as follows:
 - Page 0 is mapped to frame 2.
 - Page 1 is mapped to frame 3.
 - Page 2 is mapped to frame 4.
 - Page 3 is mapped to frame 7.
 - Page 4 is mapped to frame 8.
 - Page 5 is mapped to frame 9.
 - Frames 0, 1, 5, 6, and n are empty.

Shared Pages

- Ex.: An editor of 150K size
 - Page size: 50K
 - User data: 50K
 - 40 users → 8000K
 - Shared pages → 2150K
- Only reentrant codes can be considered as shared pages


Shared Pages



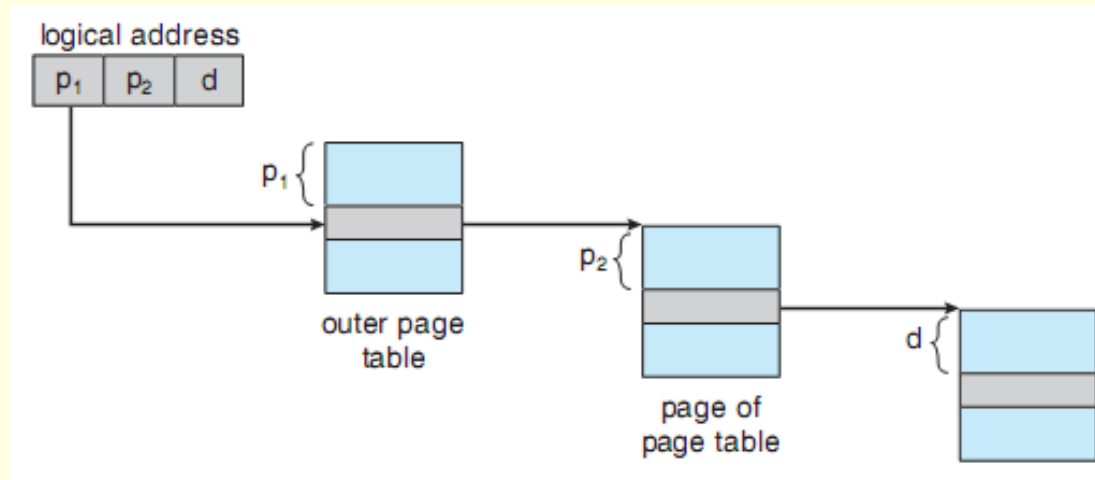

Hierarchical Paging

- Assume a system with 32 bit logical address space
- Each page is considered to be 4KB
- Page table size: 2^{20} entries
- Each entry is 4 bytes
- Page table size: 4MB
- Is it stored in contiguous memory?
- *One solution*: paging the page table

Hierarchical Paging

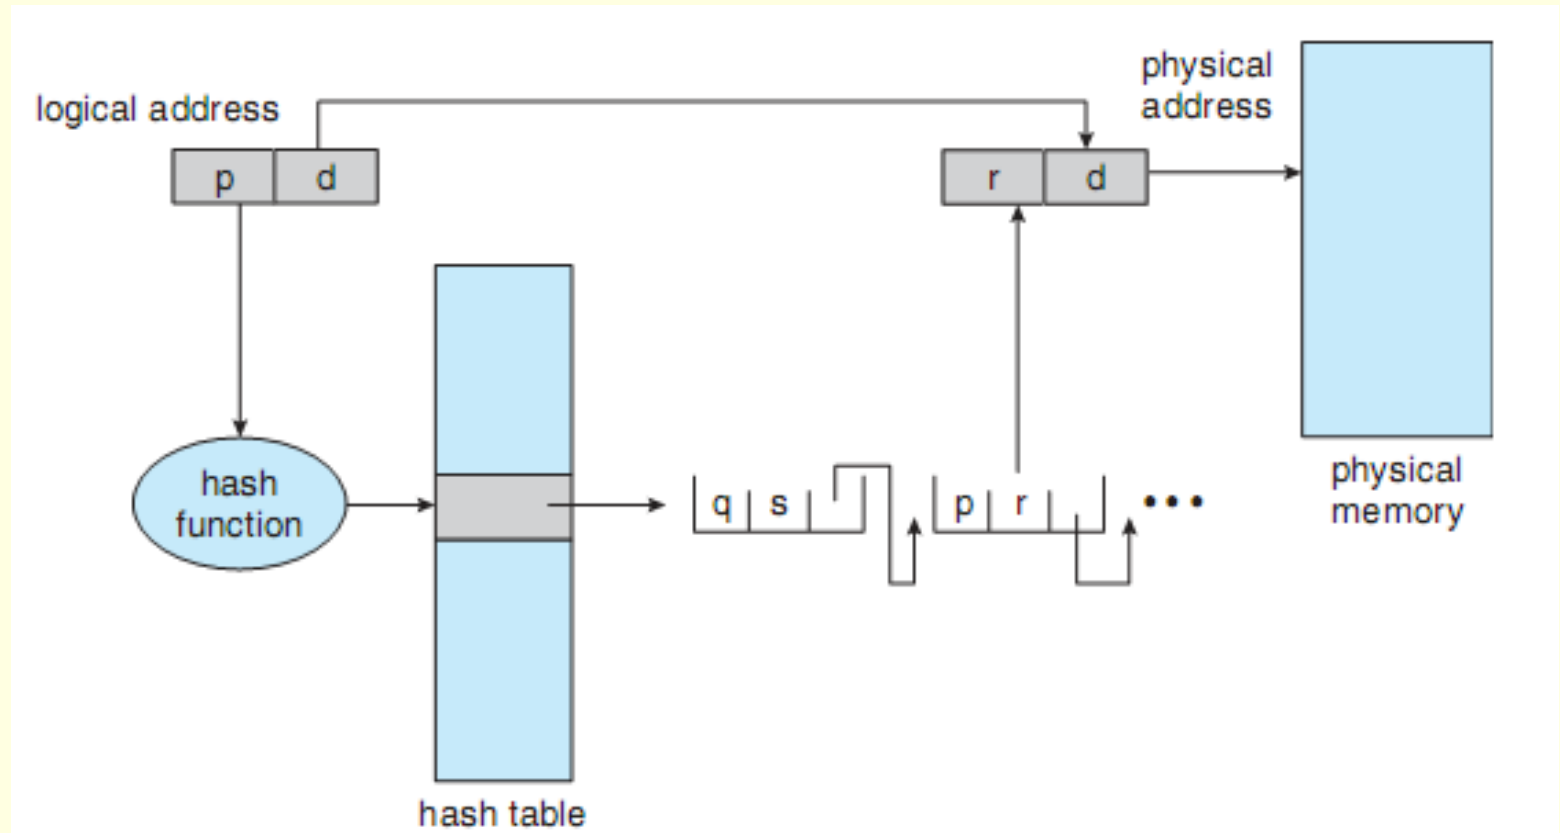


Hierarchical Paging



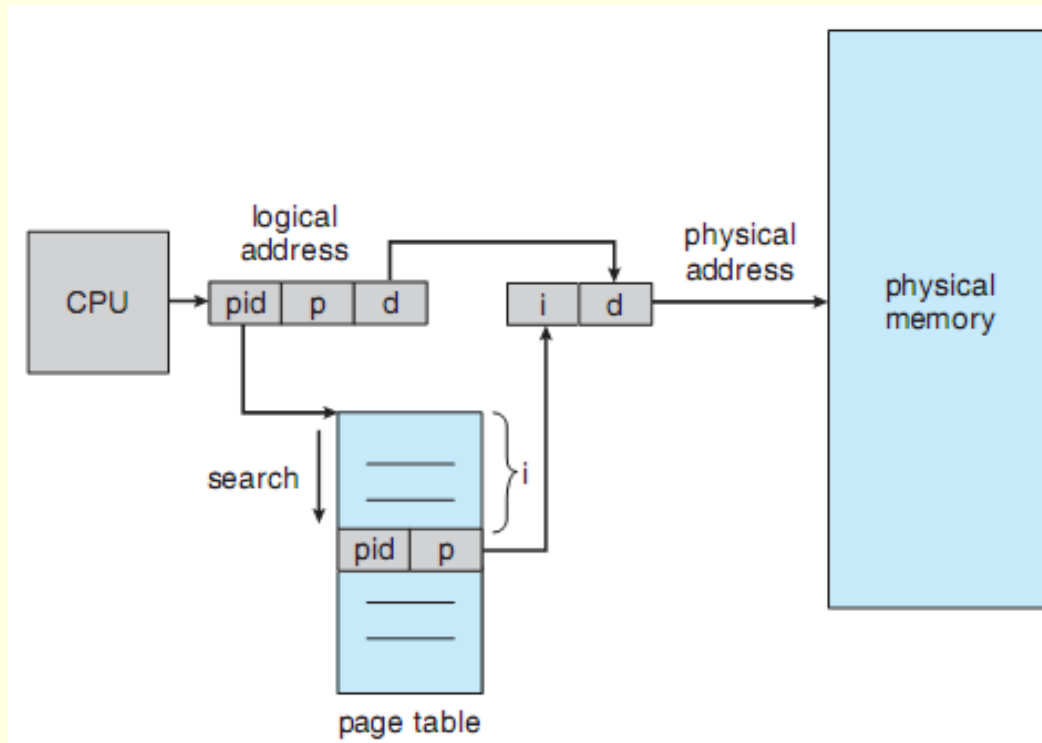
- How many level of paging is required for a 64 bit computer with 4KB pages?
 - 7 levels of paging

Hashed Page Tables



Inverted Page Tables

- Standard page tables use very large amounts of memory




- Disadvantage: Time overhead, difficulty with shared pages.

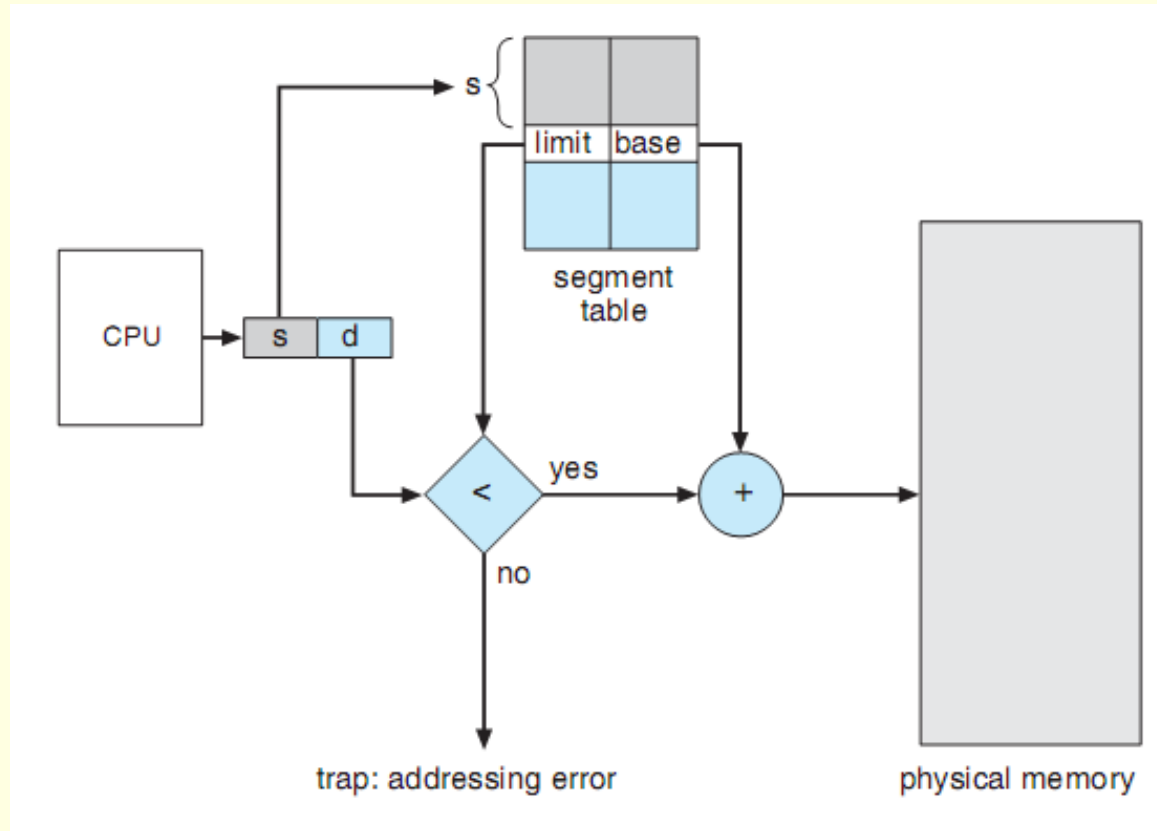
Segmentation

- User's view of memory vs. actual physical memory
- Paging: one virtual address
- Segmentation: a two-partition address
 - A segment-name (or segment-id)
 - An offset
- Some segments generated by a C compiler

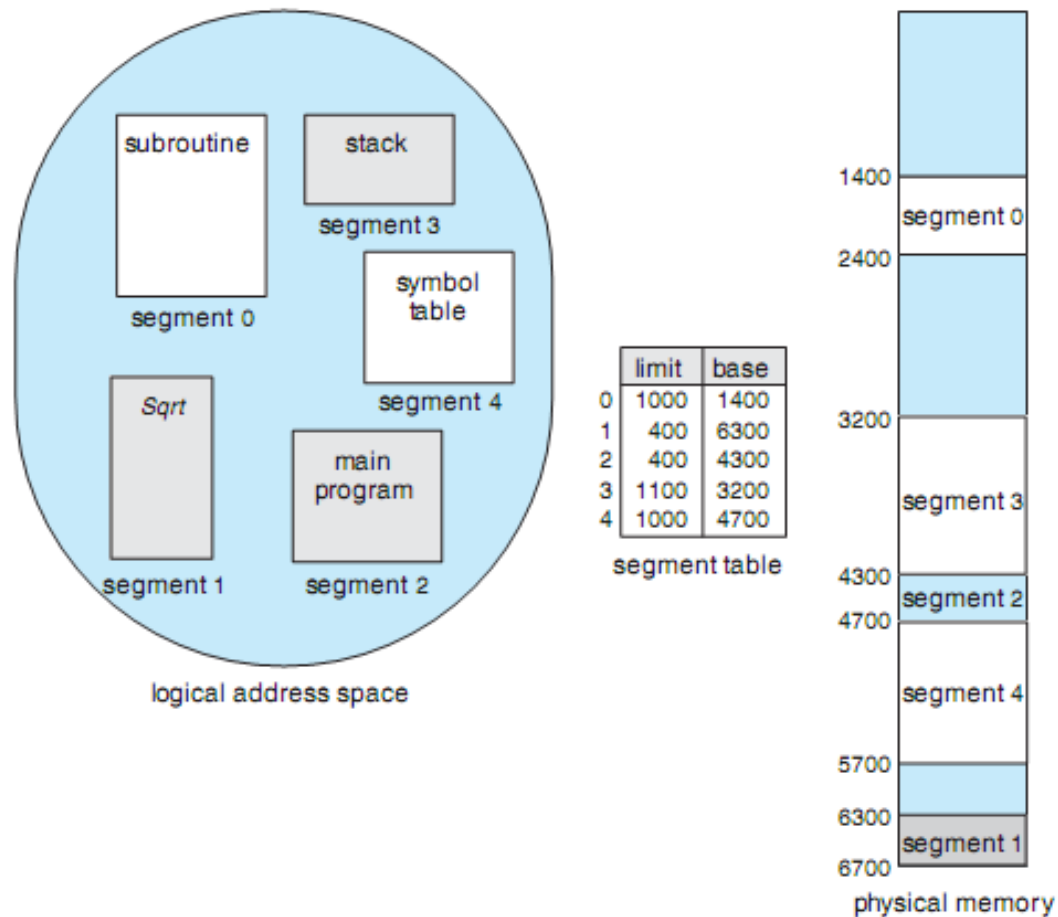
1. The code
2. Global variables
3. The heap, from which memory is allocated
4. The stacks used by each thread
5. The standard C library



Segmentation Hardware




Example



Properties

- Two memory references
- External fragmentation


Example: The Intel Pentium (Including segmentation with paging)



- Segmentation unit+ paging unit replaces MMU

- Pentium paging

- 4KB and 4MB



- Study Linux on Pentium Systems!!!

